


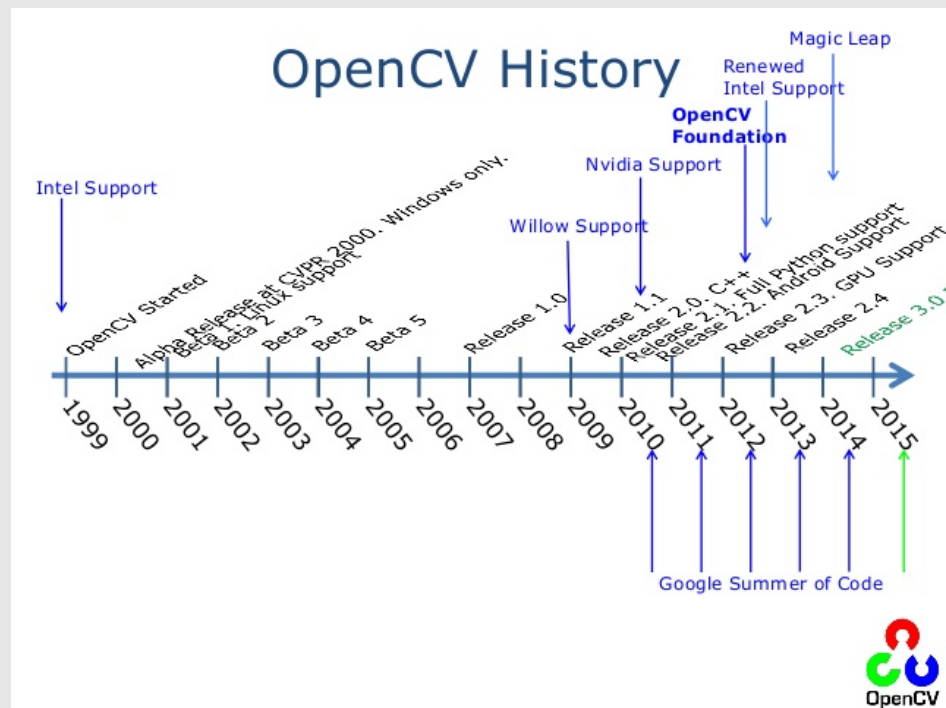
Introduction to OpenCV



OpenCV History

- OpenCV is an open library with programming functions for implementing computer vision tasks:
 - Main interface in C++
 - Also interfaces in Java, Matlab and **Python**:
<https://opencv-python-tutroals.readthedocs.io>

- History :
 - 1999 : Original version by Intel Research
 - 2006 : Version 1.0
 - 2008 : Support by Willow Garage
 - 2009 : Version 2.0 (**OpenCV 2**)
 - 2012 : Support by OpenCV.org
 - 2015 : Version 3.0 (OpenCV 3)
 - 2018 : Version 4.0 (OpenCV 4)




Source: <http://fr.slideshare.net/embeddedvision/e04-open-cvbradsk>

OpenCV-Python Installation

- Command for installation in Ubuntu/Raspbian:
sudo apt-get install libopencv-dev python-opencv python-matplotlib
- For Windows, follow tutorial in :
https://opencv-python-tutroals.readthedocs.io/en/latest/py_tutorials/py_setup/py_setup_in_windows/py_setup_in_windows.html
- For verifying, execute python in terminal :

```
>>> import cv2  
>>> print cv2.__version__
```
- Why using Python ? :
 - Programming simplicity
 - Support of library **Numpy** :
 - Optimised for numerical operations
 - Similar syntax to Matlab



OpenCV Functions

HighGUI:
I/O, Interface



Image Processing



Transforms




Fitting




Optical Flow
Tracking




Segmentation



Calibration




Features
VSLAM



Depth, Pose
Normals, Planes,
3D Features



Object recognition
Machine learning



Computational
Photography

CORE:
Data structures, Matrix math, Exceptions etc


Read/Write of Image Files

- Commands for capturing images with the camera :
 - With PiCam: *raspistill -o cam.jpg*
 - With Ubuntu Webcam: *cheese*
- Reading/Writing an image (*file.py*) :

```
import numpy as np  
import cv2
```



```
img= cv2.imread('cam.jpg', cv2.IMREAD_GRAYSCALE)  
cv2.namedWindow('Image', cv2.WINDOW_NORMAL)  
print img.shape # (rows,cols,channels)  
cv2.imshow('Image', img)  
k = cv2.waitKey(0) & 0xFF  
if k == ord('s'): # 's' key  
    cv2.imwrite('camGray.png', img)  
cv2.destroyAllWindows()
```



Capturing video from PiCamera

- Install PiCamera package:


```
sudo apt-get install python-picamera
```

- Recovering video frames from PiCamera (*cameraPi.py*) :

```
from picamera.array import PiRGBArray  
from picamera import PiCamera  
import time, cv2
```

```
camera= PiCamera()  
camera.resolution= (320,240)  
rawCapture = PiRGBArray(camera, size=(320,240))  
time.sleep(2) # Allow the camera to warmup
```

```
for frame in camera.capture_continuous(rawCapture, format='bgr',  
    use_video_port=True):  
    image= frame.array # Get raw NumPy array of the frame  
    cv2.imshow('Frame',image) # Show frame  
    rawCapture.truncate(0) # Clear stream for next frame
```



Capturing video from webcam

- Recovering video frames from webcam (*camera.py*):

```
import numpy as np
import cv2
```

```
cap = cv2.VideoCapture(0)
```

```
while(True):
```

```
    # Capture frame-by-frame
```

```
    ret, frame = cap.read()
```

```
    # Our operations on the frame come here
```

```
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
```

```
    # Display the resulting frame
```

```
    cv2.imshow('frame',gray)
```

```
    if cv2.waitKey(1) & 0xFF == ord('q'):
```

```
        break
```

```
# When everything done, release the capture
```

```
cap.release()
```

```
cv2.destroyAllWindows()
```



Modify this code for discovering the size of each frame and reduce its resolution to a half.

Use :
cap.get(propId)
cap.set(propId, value)

Modify the code for showing the frames in color

Binary Operations with Images

- How to overlap one image over another (*bitwise.py*):

Load two images

```
img1 = cv2.imread('cam.jpg')
```

```
img2 = cv2.imread('opencv_logo.png')
```

I want to put logo on top-left corner, So I create a ROI

```
rows,cols,channels = img2.shape
```

```
roi = img1[0:rows, 0:cols ]
```

Now create a mask of logo and create its inverse mask also

```
img2gray = cv2.cvtColor(img2,cv2.COLOR_BGR2GRAY)
```

```
ret, mask = cv2.threshold(img2gray, 10, 255, cv2.THRESH_BINARY)
```

```
mask_inv = cv2.bitwise_not(mask)
```

Now black-out the area of logo in ROI

```
img1_bg = cv2.bitwise_and(roi,roi,mask = mask_inv)
```

Take only region of logo from logo image.

```
img2_fg = cv2.bitwise_and(img2,img2,mask = mask)
```

Put logo in ROI and modify the main image

```
dst = cv2.add(img1_bg,img2_fg)
```

```
img1[0:rows, 0:cols ] = dst
```

```
cv2.imshow('res',img1)
```



Filtering : Convolution-Smoothing

- Filtering an image with noise (*'smoothing.py'*):

```
img = cv2.imread('bike.jpeg')  
cv2.imshow('original',img)
```

2D convolution with 5x5 mean kernel

```
kernel = np.ones((5,5),np.float32)/25  
filtered2D = cv2.filter2D(img,-1,kernel)
```

Blur with 5x5 kernel (equal to 2D convolution)

```
blur = cv2.blur(img,(5,5))
```

Gaussian blur (Gaussian kernel)

```
GaussianBlur = cv2.GaussianBlur(img,(5,5),0)
```

Median blur

```
median = cv2.medianBlur(img,5)
```

$$K = \frac{1}{25} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$
$$K = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$





Original



Gaussian Noise



Salt & Pepper Noise



Test different sizes of kernel for filtering Gaussian noise in 'bike.jpg'

Test different sizes of kernel for filtering 'salt & pepper' noise in 'cameraman.jpg'

Edge Detection : Sobel, Laplacian

- Detecting the edges in an image ('*borders.py*') :

```
ddepth = cv2.CV_64F
img = cv2.imread('sudoku.jpg')
img = cv2.GaussianBlur(img,(3,3),0)
gray = cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
```

```
# Laplacian
laplacian = cv2.Laplacian(gray,ddepth)
```

```
# Sobel Gradient-X
grad_x = cv2.Sobel(gray,ddepth,1,0,ksize = 3)
# Sobel Gradient-Y
grad_y = cv2.Sobel(gray,ddepth,0,1,ksize = 3)
```

```
# Converting back from 64-bit floating-point (CV_64F)
# into original 8-bit unsigned integers (CV_8U)
abs_grad_x = cv2.convertScaleAbs(grad_x)
abs_grad_y = cv2.convertScaleAbs(grad_y)
abs_laplacian = cv2.convertScaleAbs(laplacian)
# Combining both gradients into one Sobel image
dst = cv2.addWeighted(abs_grad_x,0.5,abs_grad_y,0.5,0)
```

Laplacian = 2nd order derivatives in x and y

$$\text{dst} = \Delta \text{src} = \frac{\partial^2 \text{src}}{\partial x^2} + \frac{\partial^2 \text{src}}{\partial y^2}$$

Discrete approximation of Laplacian by the kernel :

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

Sobel = Discrete differentiation of image intensity

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * I \quad G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * I$$

Horizontal changes

Vertical changes

Total gradient approximation

$$G = \sqrt{G_x^2 + G_y^2} \quad \text{or} \quad G = |G_x| + |G_y|$$

Contrast Adjustment : Histogram Equalization

- How to compute histograms and equalize them (*'histogram.py'*) :



```
from matplotlib import pyplot as plt  
img = cv2.imread('dark.png')  
color = ('b','g','r')
```

```
for i,col in enumerate(color):  
    histr = cv2.calcHist([img],[i],None,[256],[0,256])  
    plt.plot(histr,color = col)  
    plt.xlim([0,256])
```

```
gray= cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)  
histr_gray= cv2.calcHist([gray],[0], None,[256],[0,256])  
plt.plot(histr_gray,color='k') # Black plot
```

```
# histogram equalization  
equ = cv2.equalizeHist(gray)  
histr_equ= cv2.calcHist([equ],[0], None,[256],[0,256])  
plt.plot(histr_equ,color='m') # Magenta plot  
plt.show()
```

Histogram = Frequency of pixel intensities



Histogram Equalization

Mouse Events Handling

- Drawing circles by double-clicking on the window ('*mouse.py*'):

```
import cv2
```

```
import numpy as np
```

```
def draw_circle(event, x, y, flags, param): # mouse callback
```

```
    if event == cv2.EVENT_LBUTTONDBLCLK:
```

```
        cv2.circle(img, (x,y), 50, (255, 0, 0), -1)
```

```
img = np.zeros((512,512,3), np.uint8) # black image
```

```
cv2.namedWindow('image')
```

```
cv2.setMouseCallback('image',draw_circle)
```


```
while(True):
```

```
    cv2.imshow('image',img)
```

```
    if cv2.waitKey(1) & 0xFF == 27: # ESC key
```

```
        break
```

```
cv2.destroyAllWindows()
```



Trackbar Handling

- Choosing a color with 3 RGB trackbars (*'trackbar.py'*):

```
import cv2
```

```
import numpy as np
```

```
def nothing(x):
```

```
    pass
```

```
img = np.zeros((300,512,3), np.uint8)
```

```
cv2.namedWindow('image')
```

```
cv2.createTrackbar('R','image',0,255,nothing)
```

```
cv2.createTrackbar('G','image',0,255,nothing)
```

```
cv2.createTrackbar('B','image',0,255,nothing)
```

```
while(True):
```

```
    r = cv2.getTrackbarPos('R','image')
```

```
    g = cv2.getTrackbarPos('G','image')
```

```
    b = cv2.getTrackbarPos('B','image')
```


```
    img[:] = [b,g,r]
```

```
    cv2.imshow('image',img)
```

```
    if cv2.waitKey(1) & 0xFF == 27: # ESC key
```

```
        break
```

```
cv2.destroyAllWindows()
```





ROI selection with the mouse

- Choosing a ROI with the mouse and copy it ('roi.py'):

```
import cv2
img= cv2.imread('cam.jpg')
drag=False, point1=None, point2=None, selection=False

def choose_roi(event, x, y, flags, param): # mouse callback
    global drag, point1, point2, selection
    img2=img.copy()
    if event == cv2.EVENT_LBUTTONDOWN:
        point1=(x,y)
        drag=True
    if event == cv2.EVENT_MOUSEMOVE and drag:
        point2=(x,y)
        cv2.rectangle(img2,point1,point2,(0,0,255),4)
    if event == cv2.EVENT_LBUTTONUP and drag:
        point2=(x,y)
        cv2.rectangle(img2,point1,point2,(0,0,255),4)
        drag=False
        selection=True
    cv2.imshow('image',img2)

cv2.setMouseCallback('image', choose_roi)
while(True):
    if selection:
        cv2.imshow('roi',img[point1(0):point2(0),point1(1):point2(1)])
        selection= False
```



Canny Edge Detection


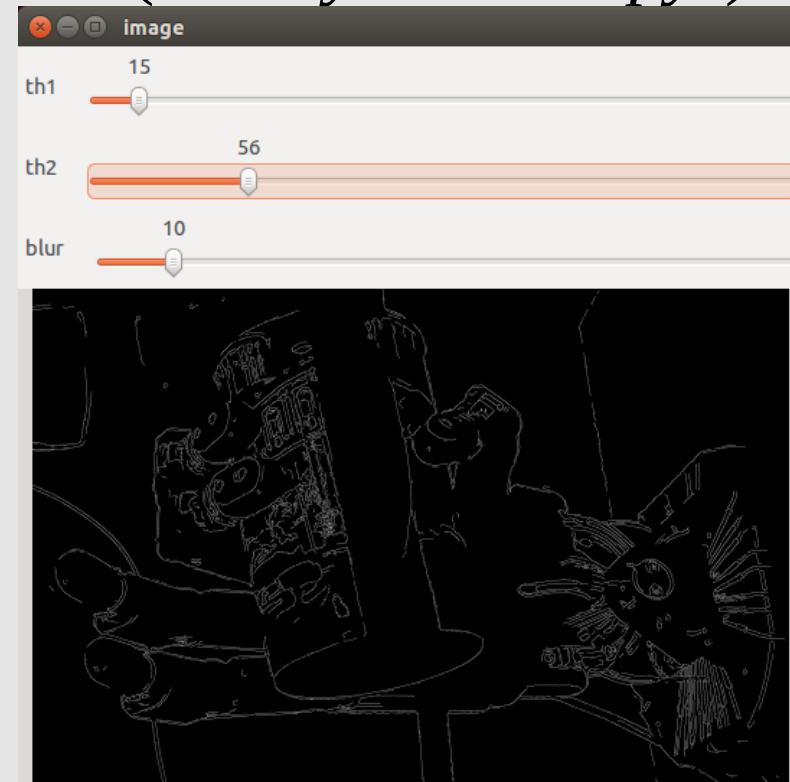
- Fixing thresholds for Canny edge detection (*'cannyDetector.py'*) :

```
import cv2
import numpy as np
```

```
img= cv2.imread('cam.jpg')
gray=cv2.cvtColor(img,cv2.COLOR_BGR2GRAY) # Change to gray
def nothing(x):
    pass
```

```
cv2.namedWindow('image',cv2.WINDOW_NORMAL)
cv2.namedWindow('blurw',cv2.WINDOW_NORMAL)
cv2.createTrackbar('th1','image',0,255,nothing)
cv2.createTrackbar('th2','image',0,255,nothing)
cv2.createTrackbar('blur','image',1,10,nothing)
```

```
while(True):
    th1=cv2.getTrackbarPos('th1','image')
    th2=cv2.getTrackbarPos('th2','image')
    size=cv2.getTrackbarPos('blur','image')
    kernel= np.ones((size,size),np.float32)/(size*size)
    filteredGray= cv2.filter2D(gray, -1, kernel) # Convolution with kernel
    edges=cv2.Canny(filteredGray, th1, th2) # Edge detection with Canny
    cv2.imshow('image',edges)
    cv2.imshow('blurw',filteredGray)
    if cv2.waitKey(1) & 0xFF==27: # ESC key
        break
cv2.destroyAllWindows()
```



Contour segmentation + moments

- Segmenting contours et getting their moments (*'moments.py'*):


```
import cv2
import numpy as np
```

$$M_{ij} = \sum_x \sum_y x^i y^j I(x, y)$$



```
def thresh_callback(thresh):
    edges = cv2.Canny(blur,thresh,thresh*3)
    drawing = np.zeros(img.shape,np.uint8)
    contours,hierarchy = cv2.findContours(edges,cv2.RETR_TREE,cv2.CHAIN_APPROX_SIMPLE)
    for cnt in contours: # For each contour
        moments = cv2.moments(cnt)
        if moments['m00']>0: # area of contour
            cx = int(moments['m10']/moments['m00']) #Centroid x
            cy = int(moments['m01']/moments['m00']) #Centroid y
            cv2.drawContours(drawing,[cnt],0,(0,255,0),1)
            cv2.circle(drawing,(cx,cy),5,(0,0,255),-1)
    cv2.imshow('output',drawing)
```

```
img = cv2.imread('cam.jpg')
gray = cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
blur = cv2.GaussianBlur(gray,(5,5),0)
cv2.namedWindow('output',cv2.WINDOW_NORMAL)
thresh = 100
max_thresh = 255
cv2.createTrackbar('canny_thresh:', 'output', thresh, max_thresh, thresh_callback)
thresh_callback(200)
if cv2.waitKey(0) == 27: #Wait until ESC
    cv2.destroyAllWindows()
```



Exercise :

Tracking an object with color

Exercise (Use file 'camera.py' or 'cameraPi.py' as base) :

- Goal : Segment an object according to its color
- Inputs : Video frames from camera + desired color
- Output : Centroid of the biggest object
- Steps :
 1. Transform image from BGR to HSV
 2. Binarize image for detecting color : `cv2.inRange`
 3. Find contours of binarized image (`cv2.findContours`)
 4. Choose biggest contour (`cv2.contourArea`), compute centroid (`moments.py`) and draw it.
- Bonus :
 - Choose binarization thresholds with trackbars
 - Choose color (H) with the mouse by clicking on it
 - Choose thresholds automatically from a ROI
 - Draw a bounding box around the object

